

Gaussian Discriminant Analysis

CS6140: Machine Learning

Uzair Ahmad

Introduction

Previously, we've discussed the Naïve Bayes Model, where components of the input vector x_i are discrete-valued. When features of x_i are continuous-valued random variables, instead of using the Gaussian Naïve Bayes Model, we could use the Gaussian Discriminant Analysis Model (GDA), in which we assume that, given the label y_i , the input x_i follows a multivariate Gaussian distribution. The GDA Model is also a generative model that can be applied to classification tasks.

GDA model parameters

GDA is also a generative model, thus we have $P(y_i|\mathbf{x}_i) = \frac{P(\mathbf{x}_i|y_i)P(y_i)}{P(\mathbf{x}_i)}$. Since the task is to classify an example, here we could make an assumption of the label such that y_i follows a Bernoulli distribution specified by parameter π . $y_i \sim \text{Bernoulli}(\pi)$.

Then, given the fact that the label is known, we can make another assumption about the input variables such that $x_i|y_i = 1$ and $x_i|y_i = 0$ follow the multivariate Gaussian distribution specified by (μ_0, Σ) and (μ_1, Σ) , respectively.

$$x_i|y_i = 0 \sim N(\mu_1, \Sigma)$$

$$x_i|y_i = 1 \sim N(\mu_0, \Sigma)$$

Given the distribution parameters are known, we have

$$P(y_i) = \pi^{y_i}(1 - \pi)^{1-y_i}$$

$$P(\mathbf{x}_i|y_i = 0) = \frac{1}{(2\pi)^{\frac{m+1}{2}}|\Sigma|^{\frac{1}{2}}} \exp\left(-\frac{1}{2}(\mathbf{x}_i - \mu_0)^T \Sigma^{-1}(\mathbf{x}_i - \mu_0)\right)$$

$$P(\mathbf{x}_i|y_i = 1) = \frac{1}{(2\pi)^{\frac{m+1}{2}}|\Sigma|^{\frac{1}{2}}} \exp\left(-\frac{1}{2}(\mathbf{x}_i - \mu_1)^T \Sigma^{-1}(\mathbf{x}_i - \mu_1)\right)$$

where $\mu_0, \mu_1 \in \mathbb{R}^m$ is the mean vector; $\Sigma \in \mathbb{R}^{m \times m}$ is the covariance matrix; and the $|\Sigma| \in \mathbb{R}$ is the determinant of Σ .

Notice that computing the probability of x under each class conditional density is equivalent to calculating the distance from x to the center of each class, μ_i , using Mahalanobis distance.

$$d = \sqrt{(x_i - \mu_1)^T \Sigma^{-1}(x_i - \mu_1)}$$

Therefore, GDA can be thought of as a nearest centroids classifier.

Training phase

According to the assumptions we've made above, GDA has the following parameters.

- π , which specifies $P(y_i)$;
- μ_0, Σ which specifies $P(\mathbf{x}_i|y_i = 0)$;
- μ_1, Σ which specifies $P(\mathbf{x}_i|y_i = 1)$;

Thus, we could write down the log-likelihood of the data

$l(\pi, \mu_0, \mu_1, \Sigma) = \ln \prod_{i=1}^N P(\mathbf{x}_i|y_i)P(y_i)$. By Maximizing l , we could get the optimal parameters as followings

$$\begin{aligned}\pi^* &= \frac{\sum_{i=1}^N \mathbb{1}(y_i=1)}{N} \\ \mu_0^* &= \frac{\sum_{i=1}^N \mathbf{x}_i \mathbb{1}(y_i=0)}{\sum_{i=1}^N \mathbb{1}(y_i=0)} \\ \mu_1^* &= \frac{\sum_{i=1}^N \mathbf{x}_i \mathbb{1}(y_i=1)}{\sum_{i=1}^N \mathbb{1}(y_i=1)} \\ \Sigma^* &= \frac{\sum_{i=1}^N (\mathbf{x}_i - \mu_{y_i})(\mathbf{x}_i - \mu_{y_i})^T}{N}\end{aligned}$$

where $\mathbb{1}(y_i = 0)$ is the indicator function such that if $y_i = 0$, the $\mathbb{1}(y_i = 0)$ outputs 1, otherwise 0.

Prediction phase

After a model is trained, we can make prediction on the label of a given data such that

$$y_i = \operatorname{argmax}_{y_i} P(\mathbf{x}_i|y_i, \theta)P(y_i).$$

If class priors are uniform, then the test data point can be classified finding

$$P(x|y_i, \theta) = \operatorname{argmin}_{y_i} P(\mathbf{x}|y_i) = (\mathbf{x} - \mu_i)^T \Sigma_{y_i}^{-1} (\mathbf{x} - \mu_i)$$

Here we assume that the covariance is common among all classes. In case each class has a different covariance, the resulting boundary will be quadratic also known as **Quadratic**

Discriminant Analysis.

```

1  class GDA():
2      def __init__(self):
3          self.pi = None
4          self.mu0 = None
5          self.mu1 = None
6          self.sigma = None
7
8      def fit(self, x, y):
9          self.pi = np.mean(y)
10         self.mu0 = np.mean(x[y[:,0]==0], axis=0)
11         # centroid of class 0
12         self.mu1 = np.mean(x[y[:,0]==1], axis=0)
13         # centroid of class 1
14
15         n_x = x[y[:,0] == 0] - self.mu0
16         p_x = x[y[:,0] == 1] - self.mu1
17
18         self.sigma = ((n_x.T).dot(n_x) + (p_x.T).dot(p_x))/x.shape[0]
19         self.sigma_inv = np.linalg.inv(self.sigma)
20
21     def predict(self, x):
22         p0 = np.sum(np.dot((x-self.mu0), self.sigma_inv)*(x-
23 self.mu0),axis=1)*self.pi
24         p1 = np.sum(np.dot((x-self.mu1), self.sigma_inv)*(x-
25 self.mu1),axis=1)*self.pi
26         return p1 >= p0

```

